

MARK STRINGER

DELIVERING THE IMPOSSIBLE



Seven metaphors for successful
software development

Delivering the Impossible

Mark Stringer

Chapter 1 - Introduction

This book is called “Delivering the impossible.” But of course. The only things that can be delivered are possible things. So I might ask you, and you might ask yourself, Why on earth did you pick up this book? It has a contradiction, an impossibility right on the front cover.

But if I did that, I would be being more difficult than you want the author of a book that promises to help with impossible things. Here are some reasons why I think you might pick up this book.

How about this? You are working on a project, maybe you’re a member of the project team, maybe you’re the manager, and you feel stuck. It seems like the project is impossible. What’s being promised can’t be delivered, nobody really knows what’s being promised, or what’s being promised can’t be delivered on time.

Every direction that you look in there seem to be problems that can’t be solved. You have found this book. You think it might be worth a read.

You have a group of people working on a project for you, and you’ve got this uncomfortable feeling that the project really isn’t going well and, for the sake of the company or for the sake of your career, you really need this project to succeed. You’d wonder if there any gems of insight in this book that you could pass on to the team or maybe anything that *you* could do differently.

Even though we think the project is impossible, you haven’t just given up. And it’s worth asking yourself why? If the project seems so impossible, why hasn’t it be stopped, abandoned, killed? Put to sleep. Why would anybody carry on with it?

Well, part of the reason is your inherent optimism. You know that a lot of things that seem impossible, are actually possible, give the right skills and expertise. You know from other parts of your life, that if you can only learn to look at things in the right way and do the right things at the right time, things that seem impossible can magically be made to work out.

So what you would really like and what you might expect to find in a book with a title like “Delivering the Impossible” is some kind of guide, some kind of handbook for dealing with projects that seem impossible.

In such a book, you might expect to find methods for spotting issues that might make a project “impossible” as early as possible. You might then expect the book to go on to talk about what to do in these situations that seem impossible. How to make a project reasonable and deliverable and give it a chance of success.

Finally, you might not be surprised if such a book also talked about the things you might look for to satisfy yourself that a project is genuinely impossible and gave advice on what to do in such circumstances, to keep ourselves and the people around us financially, physically and psychologically safe, so that you can

walk away from a lost cause and move on to succeed on other projects, ride into the sunset to fight another day.

If you're getting really demanding as a reader, you might also expect the book to talk about how deliver, not just on impossible projects, but also on impossible programmes, even in impossible organisations. Yes, if you're really demanding, you might want such a book to talk about how to deliver the impossible at scale.

Well, this is that book. And it tries to meet all of these expectations. This is especially that book if your project, programme or organisation is in some part dependent on the development of software.

If you read this book, you will definitely learn, how to quickly identify situations that seem impossible and how to manage and manipulate those situations so that's there's a chance, not just of delivery, but also of success. You'll also get lots of advice about how to look after yourself and those around you while you're doing it.

Chapter 2 - Agreed activity

One of the great things about living in London is that you can sign up to all sorts of classes. Over the past ten years, I've repeatedly signed up for improvisation classes. Why? Because my experience is that every time I go to an improvisation class, I learn something new. More specifically, I have new experience, and these often, in some way or other, turn out to be useful in real life.

One of the key ideas in improvisation is blocking. For example. If an improvised scene starts with someone knocking on a door. The scene might go something like this.

Person 1: [Knocks]

Person 2: [Opens imaginary door]

Person 1: Hello! I've brought you a cabbage.

OK, now we've reached a key stage in the drama. Because Person 2 has lots of options. Person 2 could be really happy that Person 1 had brought them a cabbage.

Person 1: Oh you remembered that this is the week that I make all my Kimchi! Oh darling! You're so thoughtful.

Or they can be really angry that Person 2 has brought a cabbage.

Person 1: And that's dinner is it? You know cabbage doesn't agree with me! Remember what happened last time. You're such an ass-hole Kevin.

The thing about both of these responses is that they move the action forward. They take someone's first idea, no matter how humble, and move it forward.

In improvisation, this is known as “Yes and.” Accepting whatever your partner gives you and amplifying it. The opposite of “Yes and” is called blocking.

So for example, if we go back to what will forever now be known as the “Cabbage scene.” and get Person 1 to knock again on the imaginary door. Person 2 could respond to the “offer” (as it’s called) of the cabbage with some kind of bizarre argument.

Person 1: Hello! I’ve brought you a cabbage.

Person 2: No you haven’t that’s a Pomeranian poodle.

or alternatively:

Person 1: Hello! I’ve brought you a cabbage.

Person 2: Humph. Fine.

Generally in improvisation these kinds of responses are not recommended? Why because they “block” the scene moving forward. They don’t build on it and move it forward. The poodle response is completely trying to stop the “offer” of the cabbage. It ruins any chance of an interesting cabbage story. The story of a couple who learn to fly on the wings of their own farts. Or perhaps a story of the dangers of living with a pot of Kimchi that becomes self-aware.

The second answer is “Humph. Fine.” This is perhaps even more dangerous to a good story. It just goes nowhere. Again the story is stopped dead in its tracks. What’s really happening in such a situation is that the person who is speaking is scared. They feel out of control, so they’re doing that absolute minimum possible. The result is very boring to watch, not very nice to their fellow performers and kind of defeats the whole object of improv. But they feel safer than if they had accepted the offer of cabbage and allowed the story to progress.

Here’s another way of saying “Humph, fine.”

Imagine that instead of just two performers on stage for an improvised scene, there are a group, let’s say seven, give or take a few. As is often the way with improvised performances, the actors get a suggestion for a setting for the scene from the audience and get “The deck of a tall sailing ship.”

Here’s what might happen.

Person 1: [Putting mimed fake telescope to their eye] Look over there! On the horizon? There’s a ship, and a flag! Is that a skull and cross bones?

As with the first scene that we talked about, the other performers have a decision to make. One obvious decision here is to totally “yes and” the Pirate ship.

Person 2: Oh my God the Pirates are coming. Haul up the sails, let’s try to outrun him.

Person 3: Oh my God the Pirates are coming. The cannons! Load the cannons!

Person 4: Oh my God the Pirate are coming. Quick hide in the lifeboats.

And all of these are good selections (as long as trying to hide in the life boats doesn't actually work and result in comedy).

But what can often happen in such a situation is that instead of saying one of the things above, that might move the story forward, someone says something like this:

Person 5: Hey! Let's wash the decks!

Persons 6, 7, 8, 9 and 10: Yes! Let's!

For some of the performers on the stage, this can seem like an attractive option. But of course, for the audience, it's a very bad idea. Once the Pirate ship has been mentioned, they want to see the pirate ship arrive. They want to see what a chase between this ship and a pirate ship looks like? They want to see if this crew really can load a cannon, point it in the right direction and fire it. They want to see the pirate king and see what he'll do to the crew members who are hiding in the lifeboats.

They don't want to be treated to the sight of everyone on the stage miming mopping the floor.

But for inexperienced improvisers, the temptation to agree to wash the decks rather than wrestle with the implications of the pirate ship which is coming closer and closer, is extremely tempting.

Why? In his books "Impro" and "Impro for storytellers" Keith Johnstone describes this kind of behaviour as "agreed activity." Putting people on a stage in front of an audience makes them scared. And scared people try, instinctively to make themselves safe, even if what they're doing individually is making what the group are trying to do together - entertain people - less likely to succeed. Part of what constitutes the skill of being a good improviser is knowing how to avoid this basic need to stay safe and instead move the story forward. Accept that there's a pirate ship on the horizon and deal with what that means.

What's fascinating about this kind of behaviour is that everybody, I mean all the improvisers on stage, seem to instinctively know that it's the right thing to do. There isn't a long discussion. It happens in a second. This is in spite of the fact that, from the point of view of the audience, it's exactly the wrong thing to do. What this shows is how good we are at shying away from things that we think might be dangerous. We instinctively avoid anything that will make us have to change our behaviour or our thinking.

One way that Johnstone suggests to help the improvisers avoid "agreed activity" is to have a director who is watching the improvisers and can intervene during live shows. In this kind of set up, it's the job of the director to spot which actions and suggestions by the improvisers will move the story forward. In desperate situations she might also make suggestions herself. Imagine that a director were watching a scene where there'd been a suggestion of washing the decks. She might allow the crew 10 seconds of deck washing without advancing

the action. But then she might shout “The deck’s clean! The pirate ship is getting nearer and nearer! Deal with the Pirate ship!”

Why am I telling you this? Because of course, this idea of concrete practice has usefulness in the context of project management. As with improvisers on stage, there’s a temptation among people who work on a project to find some agreed activity. Just like improvisers, they do this as a way of avoiding having to think, or be changed. Instinctively they shy away from the realities of the problem that they’re solving.

Of course, if we pick up on what we talked about in the introduction. If we want to deliver the impossible, the first thing that we have to do is stop ignoring the bits that look like they might be impossible.

In iterative, Agile, ways of doing project management there are lots of opportunities for the team and the team manager to spot the pirate ship. Typically, every day there’s a “stand-up” meeting where the team talk about what they did the day before, what they’re going to do today, and, most importantly, if there’s anything blocking their progress.

Of course, it’s in the nature of agreed activity that if there’s one big thing that’s blocking progress, that might be the very things that no-one mentions at all in the stand up meeting. So the signs, that there is a pirate ship that needs tackling might very well not be the sight of a Jolly Roger and the sound of “Arrrrr!” Rather it might well be reports of agreed activity. Very often, the same report with mild variations for days and days.

Another sign of agreed activity and pirate ship avoidance is when a task or “Story” as they’re often called in Agile, gets planned and then just doesn’t get done. Nobody on the team decides to pick it up. It may well be that no discussion has been had within the team about why they don’t want to pick up this story and tackle it. As with the crew on the improvisational stage, they just agreed, possibly tacitly, not to tackle it. Of course, in these situations, it’s the job of the project manager to hold up with story and ask what it is about this story that means everyone is avoiding it.

So, one place to spot the pirate ship is stand-up. Another is the retrospective. That’s a meeting that happens, typically every two weeks where the team talks about how the previous sprint went. What went well, what didn’t go so well, what could be done better.

One of the strange things about project management is that very often everybody who is working on a project knows what’s wrong with the project. If they’re not saying, it’s because they’ve tried to say or have said and have been either ignored, scolded, threatened or even disciplined for pointing it out. This is the same as if, in an improv show, one of the performers shouted “Look! A pirate ship!” and a director sitting next to the stage shouting “Shut up about the pirate ship! Clean the decks!”

A leader doesn’t need to tell the team to shut up about the pirate ships and clean

the decks many times. Once they've done it once or twice they will have already caused one of the strangest recurring problems that I've seen. And I've seen it repeatedly: highly skilled, highly paid team of professionals who've completely given up trying to think for themselves.

No matter how many times they've been beaten down, a team will still respond to the right kind of opportunity to feed back. If you run retrospectives and give people an opportunity to speak, they will speak. They will talk if they think you are listening. They will notice see and appreciate you doing your best to raise their problems with the people who need to know about them and can fix them. And sooner or later, they'll mention the pirate ship. When they do, it might be difficult to keep your jaw off the floor.

For example during the retrospective for one project that I worked on, it emerged that all of the requirements for the project were coming from a business analyst. They made no sense. One of the main struggles that the development team had on any day, was making sense of what the business analyst had put in the stories. What struck me straight away was that the business analyst wasn't in the retrospective. In an agile project, the business analyst is part of the team - they're not supposed to be outside throwing in requirements. Especially if those requirements turn out to be problematic.

So my response to this problem being raised by the team was to suggest that they business analyst should be invited to stand-up meetings, and planning meetings. And retrospectives and show and tells. This was where the developers looked at each other sideways - this is always a sign that you're finally getting to the real issue.

"We're not allowed to talk to the business analyst."

"What? That's crazy! I'm sure that can't be right."

Yes it was crazy. Yes, it was right. I tried to talk to the business analyst. I emailed him asking for a meeting. I got a phone call from his boss saying I wasn't allowed to talk to him.

And of course, there was no way that the project was going to succeed in its current state. The team needed to talk to the person who could tell them what the product was supposed to do.

Side note: This project was crazier than it even sounds. It turns out there were two groups of business analysts. Technical business analysts and business business analysts. The "business" business analysts were actually allowed to talk to the "business" - the people who wanted the product. The technical team was only allowed to talk to the technical business analysts. And these two groups of business analysts - technical and business - also weren't allowed to talk to each other. They were only supposed to communicate via emailed word documents. A substantial part of making this project possible was to make sure that we had just one business analyst. This person was allowed both to talk to the people

who wanted the software and the people who were building the software. Things got a lot better after that.

So we've talked about stand up meetings and we've talked about retrospectives. There's one more Agile meeting where the "Pirate Ships" that are on the horizon might be detected. And that's the "Show and Tell." The "Show and Tell" is a meeting where the development team show the outputs of what they've been working on in the last "Sprint" - typically two weeks. Who do they show it to? The product owner who is supposed to act as a representative of all the people who want the project to happen and any other interested stakeholders who want to come along.

After the minor victory of getting the business analyst to sit with the team, I then suggested that we start to have show and tell meetings. Fraud was an issue that affected people who worked for the bank all over the world. So for the first phone conference I could hear just from the accents the geographical spread of interest. Northern Irish accents, Scottish accents, Cockney-sounding Southend accents and Indian accents from the offshore call centres in India.

In that week, the team had been doing some work on the user interface for one of the very early screens in one of the fraud detection journeys.

One of the developers put up the login screen and then clicked through the first screen and started to talk it through. Even though most people who were on the conference call were on mute, I thought I detected some kind of change in the silence. Finally someone on the line with a cockney accent said: "Erm, I thought we'd agreed that we were going to arrange cases by account name rather than by account number?" The developer who was demoing the screen looked blank. The business analyst who we'd only just set eyes on and had only just joined the team started to look worried. "No, it's in the requirements that cases should be arranged by account number." "But that makes no sense" someone with a Northern Irish accent joined in. "It's people who are victims of fraud, not just individual accounts" added someone with a Glaswegian accent. "We need to see all the accounts that someone holds, and see the activity across all of them for this to make any sense," said someone with an accent I didn't recognise.

Yes, that's right. The structure of this project was wrong from the very beginning. This problem hadn't been detected through months and months of analysis. But it took just five minutes of putting the working software (OK, I'll admit, it was only a front end) in front of the people who might use it had found the problem.

If you do stand-ups every working day with your team, if you do retrospectives and show and tells every sprint, you will find out what the problems are. Even if you do these things *badly* - you will still find out what the problems are. Not having stand-ups, not having retrospectives and show and tells is just like shouting "Scrub the decks, don't look at the pirate ship!" at your team.

So, what does understanding this notion of "agreed activity" mean for our overall aim of delivering the impossible?

Well, it means something very basic. If you work to discover the problems that your project is facing you will find them. If you then work to solve those problems, you may well be able to transform projects that seem impossible into projects that are possible. This gives you the best chance possible of actually delivering them.

This seems so obvious as to be laughable. Why then, in project after project have I found teams that aren't articulating their problems and aren't tackling them? In improv terms rather than looking at the pirate ship and doing what needs to be done when it arrives, they are washing the decks and in the process are making possible projects impossible.

Why? Because problems are scary. Problems are humbling. Problems cause conflict. Professional people who are hired to do a job are supposed to be able to do it aren't they? What does it mean if they openly admit that there are parts of the job that they can't do? Maybe it means that someone hired the wrong people.

Highlighting the problems that a project unearths can be threatening to the sponsors of a project. What if the problem that you unearth is something that they haven't thought of and they don't know how to fix? They will be tempted to avoid addressing it, possibly by attacking or threatening the team for even daring to raise the issue.

Too, too often, when someone on the stage shouts "Look a pirate ship," It's the director off stage that shouts something like "You're wrong, it's not there." Other, equally toxic variants of this are "I'm tired of this negativity," or "Maybe you're not up to the job if you think that's a pirate ship."

And that's why people don't raise problems and instead just wash the decks.

Oh dear. Well. There is a good chance that you picked this book up because you wanted to deliver what seems to be an impossible project. I'm delivering on that promise. I'm not delivering on the promise that delivering these projects is a walk in the park or a day at the beach.

Dealing with the problems that you find and making them clear to your team and your sponsors can be a rough ride. It can be hard to stick to it. It needs an ability to talk about difficult issues. It also needs the ability to be understanding when people call you not very nice names. But there there are two very good reasons why you still should do it.

Reason number one is that, despite the resistance you might encounter, solving these problems is still the best chance you've got of delivering this project that seems to be impossible. That's a really solid reason.

But reason number two is possibly for me, just as motivating: I know what happens when you don't.

Chapter 3 - Trench Warfare

Two members of my team were supposed to be putting their software on a clients machine. They'd written a script to run on a clean machine that would put our software where it needed to be. They'd tried the script on a practice machine. It had run without any problems. They kept trying to run the script on the client's machine. It kept failing.

They ran some diagnostic tests. The machines that they were putting the software on were supposed to be clean. But running the tools we had to check for problems, it turned out that they already had other software running on them. When they pointed this out to the people who were supplying the servers, they began to act very strangely.

It was a Sunday afternoon. And I was in an emergency meeting. The topic of the emergency meeting had been what we were going to do about the failure of my team to load our software on the clients machine. The client was implying very strongly that the reason we'd failed was because our guys weren't up to the job. Just like the guys in the server room, when I brought up the subject of other software already running on the servers he became evasive. He moved the topic of the emergency meeting onto what we should do about the fact that our three o'clock emergency meeting had overrun into our four o'clock emergency meeting.

Ultimately the mystery of server software already running was solved. The company that were providing the servers were in some kind of dispute with the organisation that wanted the project. While they were in dispute, they were working to the letter of their written agreement - and so providing servers that already had some software loaded, even though they knew it would break ours. That's part of why they were being so weird. But the other part is that the written agreement to provide the servers was officially secret. So even as they saw our guys failing, and they knew why, they felt they weren't legally allowed to tell us.

Eventually we found out. Eventually it was admitted that there was no chance that the project could go live that week. I left the project the next week and never heard of it again, I'm certain it never delivered.

I'm going to use the term "trench warfare" in this chapter. And I'm going to use it to mean any project which is almost definitely going to fail unless it is substantially reorganised and restructured. Why are we talking about trench warfare now? Because in the last chapter, we were talking about pirate ships. Yes, I know, I know, I'm mixing metaphors.

What I mean by trench warfare is any project that is hugely expensive in terms of money, effort and human suffering and doesn't get anywhere. A main cause of trench warfare is a refusal (or active prevention) earlier in the project to deal with the problems that it faces. The main effect, besides project failure, is that there is ever more demand for the "agreed" activity which is an alternative to

facing the problems to become ever more frantic, pointless and damaging.

Working late. Working weekends, holidays being cancelled or frowned on, “a positive, can do attitude” are signs that you’re working on a trench warfare project. As are high turnover of staff, staff being signed off long-term sick, or with stress.

One of the things that personally I’ve found most upsetting about working on trench warfare projects is when senior management say that the staff don’t look stressed enough.

A trench warfare project is a bad situation. Experience of working on one can be one of the best motivations for getting over your reluctance to tackle issues as soon as they arise. Once you’ve worked on a trench warfare project, you might be much keener to call out problems the minute you see them on your next project. But every now and then you are going to find yourselves working on a project which is working towards a ridiculous, undeliverable deadline. At the same time this project will be struggling with several other problems. This is a very common kind of undeliverable project.

I don’t know if the real life Jerry Springer actually said this or not. But in “Jerry Springer the Opera” the character of Jerry Springer says something that has stuck with me ever since: “I don’t solve people’s problems, I televise them.”

If you’re working on a trench warfare project, that’s exactly what you should be doing. Sure, go one better than Jerry, and solve the problems that you can solve, but the problems that you can’t solve? Broadcast them. Make sure everybody on the project knows what those problems are.

Why? Because someone else may, if they really want to the project to succeed, be able to solve them. Another way to look at this is to look at some of the people who are giving you and your team a hard time. Those people who are saying that your team don’t look worried or stressed enough. Those people who are saying that your team don’t really care about the project because they took the whole weekend off. It might be that some of those people are only doing those things because they have no idea what else they can do. They want to make the project more successful, they just don’t know how. At least by making clear what the problems are that are stopping you delivering, you’re giving those people more chance of to help the project if they really do want to help.

The second thing that you can do is to deliver something to some of the people who want it. This might not be anywhere near everything that’s been asked for by the deadline. But if you can find one bit of software that you can release to some people who really want it, you can completely change the nature of a project. Pull from outside a team that comes from demand from real users is a truly wonderful thing.

Thirdly, one of the most revolutionary things that you can do, is to track and show the actual progress. This might be tough, but few things are as powerful

as showing actual progress relative to the project's needs and expectations. In a way this is just a more extreme version of the "Jerry Springer" principle. This is one of the most powerful manoeuvres that I think an Agile project manager has in their toolbox. We'll go through it in detail, well at least one way to do it, in a later chapter.

Demonstrating that a project can't possibly be delivered by a deadline can result quickly in some sensible discussions about reduction of scope. This can then, sometimes, result in the possibility of delivering a small bit of working software to some of the people that want it. It may well seem that a project is dug-in, going nowhere and doing nothing but damage to the members of its team. Still, by pushing on all of these three "fronts" and getting some kind of virtuous circle going, it might be that something good can come out of even the worst trench warfare project.

It might be that something could be done to save a trench warfare project. But equally, it's important to remember that some projects simply cannot be saved. I don't think I've ever worked on a project that was being deliberately targeted to fail, but I have heard talk of such projects.

It is quite important to understand that you don't need to stay on any project that is failing. The people in the real trenches had to stay. If they didn't they'd be shot for desertion. You don't have to. Paul Simon points out that there are multiple ways to leave a relationship. There might not be as many ways of getting off project, but there are certainly a good few.

Of course, you can just find a new job. When you start to look you might be astounded to find that you're not the only person on that project who is looking. Recruiters have a very good idea about which projects are going badly.

Chapter 4 - Flowers and Fruit

This is a distinction that I picked up from reading about Taoism. Put bluntly, there are two ways that you can get paid.

You can get paid for looking good and you can get paid for doing good. In Taoist terms. Getting paid for looking good is "Flowers" and getting paid for doing good is "Fruit."

Don't get me wrong. It might seem as we go through this chapter that I'm down on flowers. I'm not down on flowers. Flowers are great. In project management, the part of flowers is played by ideas.

People have ideas, those ideas get money and that's how projects are born. Something that some ideas and most flowers have in common is that people like them straight away. They don't have to think about it. It doesn't take any thought to like a rose. Of course, some people might not like roses. But the people who do? They don't have to take time to work it out.

Something similar seems to happen with some ideas, the kind of ideas that get money to turn into projects seem to have a common structure that makes people like them without much thought. Just like flowers, those who are exposed to them tend to like them without doing much or any thinking.

Often this is in terms of easy, fast, cheap, same and all.

For example - “Does everything that the old system does, but cheaper and faster.”

“Deals with all customer enquiries, without the need for human involvement.”

“One stop shop for everybody who is trying to do this job.”

Just like flowers, ideas get paid for being attractive to people before they even think about it.

What about fruit? What’s the difference between fruit and flowers? It isn’t that they difficult to grow. There’s lots of effort involved in growing both. And although it might *look* easy to put together an attractive idea that can gets money. Clearly it isn’t, otherwise everybody would have managed to get money for their project.

No, the difference between fruit and flowers isn’t that one is easier to do than the other. The difference is that the fruit has to actually be eaten. The users of fruit interact with it in a completely different way to the users of flowers.

Users eat fruit. They only look at flowers. Nobody ever died from looking at a flower. A sour apple can give you bad indigestion and some fruit is actually poisonous.

Similarly, implemented, delivered projects have to actually give value to their users. This involves an interaction with users which is different from buying and selling flowers. Users have to actually get some good, some value out of using a software product. What ‘good’ that value, is, is radically different for different kinds of products. For example, it’s very different for a social media product than it is for a government form. But like fruit, the process of “using” software for users is one where, if it doesn’t taste right, they will spit it out. Also like fruit, if software isn’t looked after properly, it could do them harm.

To be honest, I wish I were able to cite some other reference in support of this idea of “flowers and fruit” than a 2000 year old religious text.

Because for me, moving a project from its “flowers” state to its “fruit state” is one of the most important things that anybody is doing when they are trying to deliver a project. It’s also one of the hardest, and it’s one that hardly anybody talks about.

What makes this so hard? Well, one of the things that makes it so hard is the nature of project ideas that we talked about. What makes project ideas attractive - the kinds of things that get money - is that what makes them liked doesn’t require conscious thought. Often what makes people like things straight

away is that they don't require much thought, but they do suggest some of these features: easy, fast, cheap, same and all.

The problems start when anyone tries to implement the idea. This is especially a problem if someone tries to implement the idea using an iterative framework like Scrum. In Scrum the team meet at the beginning of every work day and talk about what they're going to do that day and anything that's stopping from doing what they want to do. They do this every working day. This is called the stand up.

Once a fortnight, the team demonstrates the work that it's been doing in the last two weeks. The audience is made up of people in the organisation and also sometimes also external users who are interested. This is called the show and tell.

The team also has a private meeting where they share what went well, what didn't go so well and what they might try differently.

Can you see yet how problems might arise for the "Flowers" aspect of a project? This idea that people liked without thinking too much about it? The minute that a team starts to implement the idea of a project they will encounter problems. And if the team is using an Agile way of doing things, like Scrum, those problems will start to be reported. Right away. On day one.

So? What's the problem with problems? The problem with problems is that any problem that the team finds is like to take the shine off the idea. Any problem that the team finds is likely to make the idea look more difficult, slower, more expensive, different and partial.

So, this is the paradox at the heart of project management and project development that I wonder people don't talk about more.

Nothing makes a project appear less attractive than actually starting to do it.

So what can be done? To solve this problem? What absolutely must not be done, as we discussed in chapters two and three, is to avoid the genuine difficulties that the project is facing. That way failure or possibly trench warfare lies.

The project is entering a vulnerable, cold phase, where it is not supported and valued for the appeal of the idea. This because we are starting to find out what is wrong with that idea. At the same time there is no support and value for the reality, because there is no reality. We don't have anything to show yet.

As someone who is trying to help this project get delivered, it is important to do two things. Firstly, it's important not to hesitate to move a project through this phase. This is what we talked about in Chapter two, when we talked about agreed activity. It's also what we talked about in Chapter three. A project can get into an awful state if the problems that agreed activity is there to avoid aren't tackled.

So, someone who wants a project to succeed shouldn't do anything to stop it moving into this phase. At the same time, you should know where you're going and what kind of things can should be done to get you out the other side of this difficult phase.

OK, here's where it gets even harder. There are two kinds of thing that you should be doing to get the project through this difficult phase and on the road to success. And here's what makes this difficult. Both of the kinds of things that you should be doing aren't often popular or well-received. They are the kinds of things that the sponsors of a project are likely to think of as either a pointless waste of time or the kind of thing that should be stopped.

What are these two things? Stakeholder research, which I'll talk about in more detail in the next chapter, and working software.

What do I mean by stakeholder research? I mean as a product team developing a good understanding of who is interested in the product and what their interest is. Notice here that I'm not saying "user research" that's because stakeholder research doesn't just include users, it include other individuals, organisations, interest groups that might be interested in the product.

I'm resisting saying things like "find out as much as you can about the people who are paying for this project" or "find out everything that you possibly can about your users." Why? Because as a seasoned and grizzled project manager I'm very careful of statements that involve "all".

But I can say this. You need a strategy for mapping the ecosystem of stakeholders. You need a strategy for investigating the needs of users. You need a way of taking what you find and using it to inform what goes into working software.

So how do we do this? How do we achieve this shift? From "Flowers", where value and support for an idea comes from the idea itself to the "Fruit" where value and support comes from a real working piece of software.

Chapter 5 - Working Software

"sufficiently advanced technology is indistinguishable from magic" Arthur C. Clarke

Concept

In the Agile manifest, working software is mentioned as the second key value. >> Working software over comprehensive documentation

It's worth asking possibly, why? What was the experience of those guys who got together to talk about "lightweight" software methodologies that meant that in the final 93 word manifest, working software took up five of those words.

I don't know for sure, but my suspicion is that they had all had experience of being involved in projects where the production of working software was delayed

for a long time while the specification was agreed.

It's important to remember that for the first forty or fifty years that software development existed, that's how people thought it should be done. Software development was called software engineering and it was thought to be an offshoot of other kinds of engineering. And in other kinds of engineering, nothing is built before the production of very carefully written plans.

But what is this second principle in the Agile manifesto saying? It's saying the engineering equivalent of "have a go at building a bridge and see how you get on."

So why? Why would people who were experienced in the software development business and had got together for the express purpose of making the way that software development happened better, why would they want something about working software to be in the manifesto.

OK, let's make this about you for a minute. Think of something that you know a lot about. There will be something. It doesn't have to be anything to do with work. But think of something that you know how to do. OK now think of some aspect of that thing that someone who wasn't an expert would think was strange about how do this thing. Why do you do that thing?

There are at least a couple of possible answers. Maybe you were taught to do this thing by whoever taught you. Maybe it's just a quirk or eccentricity that you've developed. But there's one answer which is very likely and very compelling. You do things in this way because you've seen what happens if you don't. You've learned the hard way.

And I totally think that that's the reason why "working software over detailed documentation" is in the Agile manifesto. And this is a very similar reason to one that we've discussed when we were talking about avoiding agreed activity. Why was I pushing to avoid agreed activity and work with the team to tackle whatever the difficult was that was looming on the horizon? Well, there are lots of good reasons, but one of the main ones, was that I've seen what happens when you don't do this - trench warfare.

I think it's absolutely the same reason that working software is one of the four main things that are discussed in the Agile Manifesto. The people who put it there have seen what happens if you don't push for working software. But when it comes to working software, it might also be that they've seen the good things that can come from producing working software.

Conclusion

I think there are three main reasons why any software development team should be trying to develop working software sooner rather than later.

The first reminds me of a joke.

You will never be alone, if you take with you everywhere the ingredients and equipment to make a dry martini. Because, even if you think you are completely alone and stranded on a desert island, the minute that you start to make the dry martini, someone will jump out from behind a palm tree and say “that’s not how you make a dry martini.”

Something like this is true of developing working software. The more you try to get working software loaded up on the environments where it’s really going to be used, by real users, using real data, the more likely it is that people that you didn’t even know existed are going to jump out from behind palm trees and say “you’re doing it wrong.”

In my head, I always imagine the job of getting working software out in the world on a working environment like trying to escape from a prison camp. There’s only one way to find out what all the traps are that are out there in no man’s land, and that’s to try to set them off.

I’m mixing metaphors again aren’t I? Let’s stick with the man behind a tree for a moment. Most of those things that the man who jumps out from behind a tree will mention are what are known as “non functional requirements.” The software needs to be accessible. The software needs to be secure. The software is covered by some regulator in the industry that you have never heard of.

Of course, one way to tackle all the rules and regulation man traps that your software might set off is to try to take care of them in the specification before any software gets written. But in my experience (and I suspect in the experience of the people who wrote the Agile manifesto) it’s very difficult to find out exactly what you can and can’t do without trying to do something. When you’re trying to list these requirements without a piece of working software, you’re only really dealing with “known knowns”. When you try to get some working software as far as you can through the barbed wire to the outside world, you quickly start to find about “known unknowns” (you knew there would be other security measures out there in no man’s land, but you didn’t know what they were) but also known unknowns (like the guy jumping out from behind a tree).

OK. This metaphor is in a blender. But it still applies. Why try to do working software? Because if you do a man will jump out from behind a tree and tell you why you can’t and the only way to find out what there is in no-man’s land that’s stopping you and your software from escaping is to get through it. Let’s move on to the second reason why trying to create working software is a good idea.

The second reason is this, we might call it the “there’s only one way to find out.” Reason. Sometimes it’s really important to re-state very simple things. How do you find out if you can do something? By trying to do it.

I’m teasing this out from the kind of non-functional requirements, rules and regulations reasons that might be stopping you from getting some software working. This is more basic than that. These are more basic questions that

trying to get working software answers.

Is your team capable of writing this software?

Does your team have access to the tools and resources that they need to write this software?

Does the technology that you've decided to use work?

Is the organisation that you're working for capable and willing to pay for the servers, people and set-up that you need to deploy this software?

Well, really there's only one way to find out the answers to these questions and embarrassingly, the answer might not be the one that you're hoping for. I've worked on teams where the team members don't have access to the office, I've worked on teams that don't have access to the internet. Of course these are problems that can be solved, but only once they're uncovered.

So, this is the most basic reason why a team should be trying to develop software as soon as possible, because there really is only one way to find out.

But there's a third reason. Software is magical and the way that people respond to it is magical. Put simply, people don't respond to working software in the way that they respond to feature lists or specifications.

Think about it. Think about the pieces of software that you interact with every day. Are you thinking about them in terms of lists of features? When you're using some piece of software, you are using it to do something. You have other things on your mind. A funny joke that you want to share, a report that you want to write or a podcast that you want to listen to.

And this, final reason is the main reason that it's useful, important and clever, whenever you can, to try to move a project towards delivering working software earlier rather than later. By doing this, by putting working software in the hands of the people who will use it, you start to solve the flowers vs fruit problem that we talked about in the last chapter.

Because you move the discussion and the dynamic of the project away from flowers - a discussion about the list of things that a piece of software will do to fruit - something that is in people's hands that will do something that's useful for them.

I worked on a project for an organisation that was still doing pretty much all of it's business using paper documents. This was a huge organisation, and it processed a lot of paper. The project that I was working on was trying to take just one of the processes that this organisation did and manage it using an electronic document management system.

And the project had a couple of good things going for it. Firstly it was using an Agile way of doing things. Secondly, the product owner was a former clerk of the company, she knew all the other clerks and she knew their business very well.

To start with the project had a tough time. We could get servers to put the software on. The document package we'd chosen wasn't as mature as we thought it would be. But we pushed on through some early design iterations and a load of technical difficulties until we got to one show and tell where the team finally had some working software that they could demonstrate.

It was an odd turning point in the project. Because that first demonstration was so terrible. We'd managed to pare down the minimum functionality to a view of a collection of documents and then a display of the document when its title was clicked. In that first demonstration, when we clicked on the document link, a dialog window came up that said "Do you want to veiw your document?" and when the user clicked "OK" an error message appeared.

The look on the product owner's face! At this point were were about £500,000 into a £2 Million project. And all she had to show was a misspelled-spelled but that led to an error message. It was a difficult time for her and it was a difficult time for the team.

But two weeks later it was a slightly different story. Now there was a list of documents, now when the document was clicked, the chance to open the document had a button that was spelled correctly and when it was clicked, the document was displayed.

The product owner was still looking nervous, but the lines around the eyes and forehead seemed to be slightly fewer.

Not too many show and tells after that the product owner had a question. "Can I get this on a laptop so I can show it to the clerks?"

The short answer to that question was "no" because all of this nearly half a million pounds worth of software was deployed only on developer laptops. But the product owner wanting to have a version that she could take around the country and demonstrate provided a good extra reason to negotiate with the people who'd been supposed to be giving us server space.

Once that was in place, the product owner went on the road with her laptop. The demo still wasn't much. The demo was still buggy. We still even had some spelling problems. But the demo by the product owner to her own kind went over very well. She could wave past any user interface issues, or problems and she could show her old work mates the first sight of something that could really make their lives a lot easier.

She came back with a list of problems she'd encountered using the demo, and a list of suggestions for features that had come from the clerks, but the main question that she came back with was really a demand "When will it be ready?"

From that point on, the nature of the project changed. It wasn't about delivering on a list of functions, it was about rolling out to the clerks all across the country the tiny bit of working software that the product owner had shown them and then adding to that the next obvious steps.

The other interesting thing about getting some working software in front of the people who might really use it, was that it both asked and then started to answer the other two questions that I mentioned. “Can we do this?” and “Who is going to jump out and stop us.”

To the first question, the answer at first was “No.” We didn’t have any server space where we could deploy a live service. Bean counters who were working on the project for the client had been arguing for a while that the price of providing and supporting the servers was included in the contract and they should have to either commission them, pay for them or support them. This became a much harder argument to make once people in their own company were asking for the software. Suddenly they were in the way, rather than helping the company by being tough on costs, they were stopping people who worked for their organisation from getting at something that could make their lives easier. The servers appeared, and then money for staff to support them appeared and something resembling a professional roll-out of the software started to happen.

To the second question “If we try to do this, is anybody going to jump out from behind a tree and stop me?” The answer was “yes.” In fact two people jumped out. An accessibility guy and a security guy. The accessibility guy claimed there was no way the software could be released until it met an extra set of requirements that would make it usable according to a yet another set of standards that try to make websites more accessible to people with disabilities. The client claimed that we should have known about these requirements right from the start and so we should pay for them.

And by the way. This guy was pretty much right. We should have built in accessibility right from the start. It takes no more effort in coding to make sure that a website is accessible and it actually makes the site much easier for all kinds of people who you might not think of as disabled to use your site. Do it. If in doubt, pay a blind guy to look at your site.

Out from behind another tree jumped the security guy. He said that the project should never be allowed to go live until we could prove that it was secure. Just to leave us in no doubt about his effectiveness as an obstacle, he also refused to tell us what it was we needed to change so that it would be secure. And of course, the costs of any changes we made needed to be born by us rather than the client.

As people who jump out from behind trees trying to block progress go, these two looked pretty effective. Both were telling us we couldn’t release until we did what they said, both were telling us that we had to pay to do what they said. The security guy was being even more effective at blocking us because he was also not telling us what it was we had to do.

But neither of these guys was a match for the clerks. There were a lot of clerks. They’d seen that this software would make their lives much easier. We improved the accessibility, but we went live with what we had. We submitted the software

to outside security testing. We addressed some issues, but we went live with some others still being looked at.

Putting bad, barely working software in front of real users completely changed the project.

Connections

Showing barely working software to the clerks had created this magical thing, which gets talked about when people talk about other ways of doing Agile - like Lean, which is basically Agile, but for manufacturing. What is this magical thing? It's called "pull".

You may have heard the phrase "just in time." It's a strategy for managing supply chains in all kinds of industries. One industry that has perfected this approach to managing its work is the Japanese car industry, and particularly the Toyota car company.

The Toyota company succeeded because from the start it understood one thing: it just wasn't possible to make cars in Japan in the way they were being made in the 1920's in America? Why? Because Japan's economy at the time was tiny and it went in cycles of boom and bust. So from very early on, Toyota tried match the rate that it made cars to the pattern of demand for cars that it saw from the market. Over nearly a hundred years, that process has become so sophisticated that many different kinds of cars, with the many different options that modern cars have, can roll off the production line of the same plant.

Each car is *pulled* off the production line by a specific request for a car which comes from a dealer, which ultimately comes from a customer.

Thinking about things in terms of providing value to the customer results in another important idea which comes from the way that Toyota do things - waste. In Toyota's way of thinking. Anything which money has been spent on, which isn't on its way to a customer, is waste.

If you put these two ideas together: make things in response to demands from customers and don't have anything hanging around the factory and don't do anything that isn't on its way to customers, you end up with a process that fits the labels that have been applied to it in the west - lean and "Just-in-time."

It's important to point out that making cars is *very* different from making software. It's also possible that reports and descriptions of how the Japanese, and especially Toyota, make cars, might differ radically from how it's done today - or how it's ever done.

Even so, taking these two ideas - only deliver things which the user wants and don't do anything which isn't going to either directly benefit the user and really useful. They are both really two sides of the same idea and if taken seriously, they are a really powerful way of making projects which seem impossible start to suddenly seem possible.

Concrete Practice

So how do we do this? Every project is a little bit different. But I'm going to arbitrarily invent a rule of sixths. Which feels about right to me. If you think your project is about 6 months long. I'll give you a month to do some set up. But if after a month you and your team don't have *something* that works, you've waited too long. By the end of a sixth of the time, you need tiny piece of software that starts to do the thing that you're supposed to be doing in this project that you can demonstrate.

And once you've got this tiny little thing, you should be looking to improve its status in some way. And there are lots of ways to do this. Sure, add functionality. But if this is just a prototype that's running on a laptop, that's not the first way to go. Just as in the example that I gave. If you've got an example that works on a laptop, then one good move would be to show it to users. Another would be to move what's on the laptop to test servers and beyond that, to the live servers where it will eventually be used.

What's so important with working software, is not so much what you've got, but that you've got something and it's moving in all the right directions. Sure it needs to be increasing in terms of what it can do. But it also needs to be moving from developer environments, to test environments and onto live environments. And it needs to be moving from being tested by the team to being tested by "friendly" users to being tested by complete strangers. If it's a business application, it needs to move from using dummy data, to using data that looks like live data, to using real data.

I know nothing about rock climbing. But this is a little bit like if you're climbing a huge rock face. You climb a bit, then you put in one of those things that holds the rope to the rock. Then you pull on it, to make sure it's firm and would hold you if you fell, then you can climb a bit more.

Yes, delivering software is a bit like that. It's a bit like climbing a rock face. The weird thing is how many people think that you can get to the top without a rope - but also, without the climb!

Gradually developing working software in all of those dimensions is the powerful careful, effective way to deliver something which seems impossible. And if the thing that you're doing is in any way useful or interesting to the people who will use it, at some point in that process, we hope sooner rather than later, you will start to get "pull". You will start to get demand for the software and demand for functionality from the software from the people who are likely to use it, rather than "push" from the people who got the funding for the project.

Gradually developing software in these directions will also result in people jumping out from behind trees and helpfully providing you with extra rules and regulations that you need to follow before your software can finally escape and live free in the real world. It's (one of) the most powerful secrets to delivering things that seem impossible. And everybody, or nearly everybody, will try to

stop you doing it.

What? Yes, that's right. Incrementally delivering working software is the thing that you absolutely need to do to succeed and nearly everybody will try to stop you doing it.

How? Why?

OK, let's deal with the how and the why, but separately.

How? Developers will tell you that there's no point breaking big bits of functionality into smaller bits that could be demonstrated, released and tested. They'll say that it "only makes sense" to release some bit of functionality in one big chunk.

How? Somebody will tell you that there's no point bothering users with small bits functionality which don't show the whole journey, that it's a waste of time and money.

How? Somebody possibly even some users will tell you that there's no point looking at the new system until it has their real data in it, or until when they press the "launch missile" button something actually happens.

How? Somebody will tell you that test servers are expensive.

How? Somebody will tell you that the live environment is only in the budget from the week before the project is about to finish?

How? Nearly everybody (I've been privileged to work with some exceptions) who is responsible for product management will duck their responsibility for prioritising what the software should do. They will utter the magic tragic words "we need it all, so I don't think it's important what order it gets done in."